

Data Preparation Process



- ABC Company Digital Media Data.xlsx had missing values and categorical variables that needed attention.
- After visualizing the data, the missing values were replaced with the column mean using R programming language.
- The categorical variables were dummy coded using the Pandas library in Python.
- The continuous variables were also normalized with the use of Python programming language

Modeling: Decision Tree

import library for decision tree

```
[4]: #import tlib for decision tree
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score, confusion_matrix
```

Split into features (X) and target variable (y)

```
[5]: # Split into features (X) and target variable (y)
X = filtered_df.drop('Engaged_Visits_Post_Click', axis=1)
y = filtered_df['Engaged_Visits_Post_Click']
```

split data into test and training set 80:20

```
[16]: #split data into test and training set 80:20
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

import decision tree regressor

```
[17]: from sklearn.tree import DecisionTreeRegressor
# Create and train the decision tree classifier
# Decision Tree for Classification
regressor = DecisionTreeRegressor()
regressor.fit(X, y)
```

```
[17]: ▾ DecisionTreeRegressor
DecisionTreeRegressor()
```

```
[18]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeRegressor
from sklearn.metrics import mean_squared_error, r2_score
```

Make predictions on the test set, calculate MSE on test set and R^2

```
[19]: y_pred = regressor.predict(X_test)

# Evaluate the model's performance using Mean Squared Error (MSE) and R-squared (R2) score
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print(f"Mean Squared Error (MSE): {mse}")
print(f"R-squared (R2) Score: {r2}")

Mean Squared Error (MSE): 27.11281524885022
R-squared (R2) Score: 0.9940592575037209
```

Find feature importances

```
[10]: # 1. Feature Importances
feature_importances = regressor.feature_importances_
feature_importances_df = pd.DataFrame({'Feature': X_train.columns, 'Importance': feature_importances})
feature_importances_df = feature_importances_df.sort_values(by='Importance', ascending=False)

# 2. Decision Tree Structure
tree_structure = regressor.tree_

# 3. Predictions on the test set
y_pred = regressor.predict(X_test)

# Evaluate the model's performance using Mean Squared Error (MSE) and R-squared (R2) score
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print(f"Mean Squared Error (MSE): {mse}")
print(f"R-squared (R2) Score: {r2}")

print("\nFeature Importances:")
print(feature_importances_df)
print("\nDecision Tree Structure:")
print(tree_structure)

Mean Squared Error (MSE): 27.11281524885022
R-squared (R2) Score: 0.9940592575037209

Feature Importances:
   Feature  Importance
1  Clicks      0.704363
0  Impressions 0.232591
3  Net_Cost_CM300 0.038369
2  Net_Cost_Calculated 0.023363
8  Video_Plays  0.000575
7  Video_Replays 0.000396
11 Video_5      0.000147
13 Video_25     0.000109
12 Video_Fullscreen 0.000043
9  Video_Pauses 0.000019
6  Video_Skips  0.000018
14 Video_Completions 0.000004
5  Video_75     0.000002
4  Video_Unmutes 0.000001
10 Video_Mutes  0.000001

Decision Tree Structure:
<sklearn.tree._tree.Tree object at 0x7f572f6bee20>
```

prune tree to prevent overfitting and lower R^2

Find the tree with the best max depth and random state and save that as variables best_tree

```
[22]: param_grid = {
      'max_depth': np.arange(3, 21), # Try different depth values from 3 to 20
    }

grid_search = GridSearchCV(tree, param_grid, cv=5, scoring='neg_mean_squared_error')
grid_search.fit(X_train, y_train)

best_tree = grid_search.best_estimator_
print(best_tree)

DecisionTreeRegressor(max_depth=14, random_state=42)
```

calculate R^2 and MSE of Pruned tree

```
[24]: y_pred = best_tree.predict(X_test)
mse = mean_squared_error(y_test, y_pred)
pruned_mse = mean_squared_error(y_test, y_pred_pruned)
print("Pruned Mean Squared Error (MSE):", pruned_mse)

# Make predictions on the test set
y_pred = best_tree.predict(X_test)
print("prediction")
print("\n", y_pred)
# Calculate the R-squared score
r2_pruned = r2_score(y_test, y_pred)

print("R-squared score pruned:", r2_pruned)

#R^2 is high enough for this to be considered a good model, but not too high to expect overfitting.
Pruned Mean Squared Error (MSE): 533.2064048283789
prediction
[ 3.62963141  0.59279572  0.59279572 ...  3.62963141  1.37654321
 10.39402318]
R-squared score pruned: 0.8831680915693003
```

Modeling: Gradient Boosted

Gradient Boosted Regression

Import library

```
[25]: import pandas as pd
      from sklearn.model_selection import train_test_split
      from sklearn.ensemble import GradientBoostingRegressor
      from sklearn.metrics import mean_squared_error
```

Split the data into training and testing sets

```
[26]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

Create and fit the Gradient Boosted Regressor model

```
[27]: gb_regressor = GradientBoostingRegressor()
      gb_regressor.fit(X_train, y_train)
```

```
[27]: ▼ GradientBoostingRegressor
      GradientBoostingRegressor()
```

Make predictions and evaluate the model

```
[28]: y_pred = gb_regressor.predict(X_test)
      mse = mean_squared_error(y_test, y_pred)
```

```
[29]: print("MSE for Gradient Boosted= ", mse)

MSE for Gradient Boosted= 518.9986803001185
```

```
[30]: print(f'Mean Squared Error: {mse}')

Mean Squared Error: 518.9986803001185
```

Grid search with cross-validation

```
[32]: grid_search = GridSearchCV(gb_regressor, param_grid, cv=5, scoring='neg_mean_squared_error', n_jobs=-1)
      grid_search.fit(X_train, y_train)
```

```
[32]: ► GridSearchCV
      ► estimator: GradientBoostingRegressor
        ► GradientBoostingRegressor
```

Get the best hyperparameters and model

```
[33]: best_params = grid_search.best_params_
      best_gb_model = grid_search.best_estimator_

Make predictions and calculate the MSE
```

```
[32]: grid_search = GridSearchCV(gb_regressor, param_grid, cv=5, scoring='neg_mean_squared_error', n_jobs=-1)
      grid_search.fit(X_train, y_train)
```

```
[32]: ► GridSearchCV
      ► estimator: GradientBoostingRegressor
        ► GradientBoostingRegressor
```

Get the best hyperparameters and model

```
[33]: best_params = grid_search.best_params_
      best_gb_model = grid_search.best_estimator_

Make predictions and calculate the MSE
```

```
[34]: y_pred = best_gb_model.predict(X_test)
      mse = mean_squared_error(y_test, y_pred)
      print(f'Best Hyperparameters: {best_params}')
      print(f'Mean Squared Error: {mse}')
```

Best Hyperparameters: {'max_depth': 6}
Mean Squared Error: 460.0655449488492

Modeling: Random Forest

```
In [9]: import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error, r2_score

selected_features = df.iloc[:, 1:16]
target = df['Engaged_Visits_Post_Click']

X_train, X_test, y_train, y_test = train_test_split(selected_features, target, test_size=0.2, random_state = 42)
rf_regressor = RandomForestRegressor(n_estimators=100, random_state = 42)
rf_regressor.fit(X_train, y_train)
y_pred = rf_regressor.predict(X_test)
mse = mean_squared_error(y_test, y_pred)
r_squared = r2_score(y_test, y_pred)
print("Mean Squared Error:", mse)
print("R-squared:", r_squared)

feature_importances = rf_regressor.feature_importances_
importance_df = pd.DataFrame({'Feature': selected_features.columns, 'Importance': feature_importances})
importance_df = importance_df.sort_values(by='Importance', ascending=False)
print("Feature Importances:")
print(importance_df)
```

Mean Squared Error: 389.88302433518857

R-squared: 0.9145719605291046

Random Forest Feature Importance

